

C/C++中的 `#define` 指令用于创建宏。宏本质上是一种文本替换机制，由预处理器在编译前执行。当您 `#define` 某物时，预处理器将替换定义的宏的每个出现，用其对应值或代码。

在提供的代码中，`FF`、`GG`、`HH` 和 `II` 被定义为宏，而不是函数。这就是为什么它们没有显式的 `return` 语句。它们不是返回值，而是直接使用 `+=` 和 `=` 操作符修改它们的参数（`a`、`b`、`c`、`d`）。这些更改直接应用于传递给宏的变量。这与函数有重要区别，因为函数通常在参数的副本上操作（除非你传递指针或引用）。

`FF` 宏定义了 MD5 算法中第一轮变换的步骤，其中：

- `a`, `b`, `c`, `d` 是 MD5 状态的四个 32 位字
- `x` 是输入消息的一个 32 位块
- `s` 指定左循环移位的位数
- `ac` 是一个基于正弦函数的常量

```
guess > == test.o
1    bba46eb8b53cf65d50ca54b2f8afd9db
2    原始MD5Hash结果: bba46eb8b53cf65d50ca54b2f8afd9db
3    验证结果: 相同
4
5    Authorized users only. All activities may be monitored and reported.
6
```

得到了正确的结果

效率变低了，我们需要分析原因：

内存分配和访问模式低效：

- 为每个消息单独调用 `StringProcess`，导致多次内存分配
- 从多个不连续内存位置加载数据，导致缓存利用率低

```
86    Guesses generated: 9853408
87    Guesses generated: 10106852
88    Guess time:8.23305seconds
89    Hash time:13.8043seconds
90    Train time:98.4926seconds
91
92    Authorized users only. All activities may be monitored and reported.
93
```

训练时间太久了所以我们将训练数据减少到30000

```
112    Guesses generated: 9997458
113    Guesses generated: 10097691
114    Guess time:18.7103seconds
115    Hash time:9.05261seconds
116    Train time:0.856026seconds
117
118    Authorized users only. All activities may be monitored and reported.
119
```

```
111  Guesses generated: 9897171
112  Guesses generated: 9997458
113  Guesses generated: 10097691
114  Guess time:18.6308seconds
115  Hash time:8.94697seconds
116  Train time:0.857836seconds
117
118  Authorized users only. All activities may be monitored and reported.
119
```

重新选定基准值9.05s, 8.94s, 下面是并行化之后的结果（已经优化了访存模式

```
12  Guesses generated: 9997458
13  Guesses generated: 10097691
14  Guess time:19.6023seconds
15  Hash time:13.2519seconds
16  Train time:0.846007seconds
17
18  Authorized users only. All activities may be monitored and reported.
19
```

```
Annotate MD5Hash_SIMD
Zoom into main thread
Zoom into main DSO (use the 'k' hotkey to zoom directly into the kernel)
Browse map details
Run scripts for samples of symbol [MD5Hash_SIMD]
Run scripts for all samples
Switch to another data file in PWD
Exit
```

1. Annotate MD5Hash_SIMD

- 显示MD5Hash_SIMD函数的源代码级性能注解
- 可以看到每行代码执行的频率和耗时，帮助定位具体的性能瓶颈

2. Zoom into main thread

- 只查看主线程的性能数据
- 过滤掉其他线程的数据

3. Zoom into main DSO

- 缩小视图到主要的动态共享对象
- k 快捷键可以直接查看内核相关性能数据

4. Browse map details

- 浏览内存映射的详细信息
- 查看代码和数据的内存分布

5. Run scripts for samples of symbol [MD5Hash_SIMD]

- 对MD5Hash_SIMD函数的采样数据运行自定义脚本
- 可以进行进一步的自动化分析

6. Run scripts for all samples

- 对所有采样数据运行脚本

7. Switch to another data file in PWD

- 切换到当前目录中的其他perf数据文件

8. Exit

- 退出perf report工具

Samples: 191K of event 'cycles:u', 4000 Hz, Event count (approx.): 71032229743, DSO: main
MD5Hash SIMD /home/s2313211/guess/main [Percent: local period]

Percent	Instruction
	ls1 x19, x0, #2
	mov x0, x19
	→ bl _init
	mov x2, x0
	mov x1, x2
	sub x0, x19, #0x1
4.56	17c: cmp x0, #0x0
	↓ b.lt 194
0.05	strb wzr, [x1]
	add x1, x1, #0x1
0.02	sub x0, x0, #0x1
	↑ b 17c

汇编代码中与0相比较耗费了我们较多的时间，所以我们选择打开循环，直接执行循环里面的操作

```
for (int i = 0; i < 4; i++) {
    Byte* block_ptr = paddedMessages[i] + block * 64 + j * 4;
    // 使用位运算而不是多次左移减少指令
    values[i] = (uint32_t)block_ptr[0] |
        ((uint32_t)block_ptr[1] << 8) |
        ((uint32_t)block_ptr[2] << 16) |
        ((uint32_t)block_ptr[3] << 24);
}
```

// 展开的数据加载循环

```
Byte* block_ptr0 = paddedMessages[0] + block * 64 + j * 4;
Byte* block_ptr1 = paddedMessages[1] + block * 64 + j * 4;
Byte* block_ptr2 = paddedMessages[2] + block * 64 + j * 4;
Byte* block_ptr3 = paddedMessages[3] + block * 64 + j * 4;

values[0] = (uint32_t)block_ptr0[0] | ((uint32_t)block_ptr0[1] << 8) |
    ((uint32_t)block_ptr0[2] << 16) | ((uint32_t)block_ptr0[3] << 24);
values[1] = (uint32_t)block_ptr1[0] | ((uint32_t)block_ptr1[1] << 8) |
    ((uint32_t)block_ptr1[2] << 16) | ((uint32_t)block_ptr1[3] << 24);
values[2] = (uint32_t)block_ptr2[0] | ((uint32_t)block_ptr2[1] << 8) |
    ((uint32_t)block_ptr2[2] << 16) | ((uint32_t)block_ptr2[3] << 24);
values[3] = (uint32_t)block_ptr3[0] | ((uint32_t)block_ptr3[1] << 8) |
    ((uint32_t)block_ptr3[2] << 16) | ((uint32_t)block_ptr3[3] << 24);
```

```
paddedMessages[0] = buffer + 0 * max_padded_length;
PrepareMessage(inputs[0], paddedMessages[0], &messageLengths[0]);
paddedMessages[1] = buffer + 1 * max_padded_length;
PrepareMessage(inputs[1], paddedMessages[1], &messageLengths[1]);
paddedMessages[2] = buffer + 2 * max_padded_length;
PrepareMessage(inputs[2], paddedMessages[2], &messageLengths[2]);
paddedMessages[3] = buffer + 3 * max_padded_length;
PrepareMessage(inputs[3], paddedMessages[3], &messageLengths[3]);
```

```

111  Guesses generated: 989/1/1
112  Guesses generated: 9997458
113  Guesses generated: 10097691
114  Guess time:19.6311seconds
115  Hash time:12.5921seconds
116  Train time:0.845521seconds
117
118  Authorized users only. All activities may be monitored and reported.
119

```

```

112  Guesses generated: 9997458
113  Guesses generated: 10097691
114  Guess time:19.6759seconds
115  Hash time:12.6927seconds
116  Train time:0.855539seconds
117
118  Authorized users only. All activities may be monitored and reported.
119

```

```

113  Guesses generated: 10097691
114  Guess time:19.7199seconds
115  Hash time:12.5602seconds
116  Train time:0.853786seconds
117
118  Authorized users only. All activities may be monitored and reported.
119

```

差不多优化了1秒，还是很多的

我们使用静态变量 `a0_init`，减少每次进入函数时的 `vdupq_n_u32` 运算

```

112  Guesses generated: 9997458
113  Guesses generated: 10097691
114  Guess time:19.7208seconds
115  Hash time:12.8036seconds
116  Train time:0.854116seconds
117
118  Authorized users only. All activities may be monitored and reported.
119

```

好像没什么变化

```

// 替换这样的循环：
max_length = std::max(max_length, inputs[0].length());
max_length = std::max(max_length, inputs[1].length());
max_length = std::max(max_length, inputs[2].length());
max_length = std::max(max_length, inputs[3].length());

// 使用条件移动代替条件分支
max_padding_bits = (max_padded_bits == 448) ? 512 :
    ((448 - max_padded_bits) & 0x1FF);

```

```

L1  Guesses generated: 9897171
L2  Guesses generated: 9997458
L3  Guesses generated: 10097691
L4  Guess time:19.7885seconds
L5  Hash time:12.6222seconds
L6  Train time:0.848716seconds
L7
L8  Authorized users only. All activities may be monitored and reported.

```

还是没有什么优化感觉

```
const size_t block_offsets[16] = {0, 4, 8, 12, 16, 20, 24, 28, 32, 36, 40, 44,
48, 52, 56, 60};
```

再加一个静态变量，把这个的乘法也省掉了

```
Byte* block_ptr0 = paddedMessages[0] + block * 64 + block_offsets[j];
Byte* block_ptr1 = paddedMessages[1] + block * 64 + block_offsets[j];
Byte* block_ptr2 = paddedMessages[2] + block * 64 + block_offsets[j];
Byte* block_ptr3 = paddedMessages[3] + block * 64 + block_offsets[j];
```

```

110  Guesses generated: 9796757
111  Guesses generated: 9897171
112  Guesses generated: 9997458
113  Guesses generated: 10097691
114  Guess time:19.6745seconds
115  Hash time:12.7869seconds
116  Train time:0.842985seconds
117
118  Authorized users only. All activities may be monitored and reported.
119

```

还是没变化，我们从头再来哈，这个是并行的结果19.46s

问题 输出 调试控制台 终端 端口

perf - guess

Samples: 167K of event 'cycles:u', Event count (approx.): 87800961180

Overhead	Command	Shared Object	Symbol
19.28%	main	main	[.] MD5Hash_SIMD
8.39%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_impl_data::_M_copy_data
7.75%	main	main	[.] ROTATELEFT_SIMD
4.64%	main	main	[.] std::_Vector_base<segment, std::allocator<segment> >::_Vector_impl_data::_M_copy
3.16%	main	main	[.] PT::operator=
3.01%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_impl_data::_Vector_impl_da
2.08%	main	main	[.] F_SIMD
2.04%	main	main	[.] std::vector<int, std::allocator<int> >::_M_move_assign
1.96%	main	main	[.] G_SIMD
1.86%	main	main	[.] std::_Vector_base<segment, std::allocator<segment> >::_Vector_impl_data::_Vector
1.77%	main	main	[.] H_SIMD
1.60%	main	main	[.] std::vector<int, std::allocator<int> >::operator=

Tip: Add -I to perf record to sample register values, which will be visible in perf report sample context.

13.78s

问题 输出 调试控制台 终端 端口

perf - guess

Samples: 145K of event 'cycles:u', Event count (approx.): 76020724029

Overhead	Command	Shared Object	Symbol
23.38%	main	main	[.] MD5Hash
9.44%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_impl_data::_M_copy_data
5.24%	main	main	[.] std::_Vector_base<segment, std::allocator<segment> >::_Vector_impl_data::_M_cop
3.74%	main	main	[.] PT::operator=
3.40%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_impl_data::_Vector_impl_d
2.33%	main	main	[.] std::vector<int, std::allocator<int> >::_M_move_assign
2.08%	main	main	[.] std::_Vector_base<segment, std::allocator<segment> >::_Vector_impl_data::_Vecto
1.80%	main	main	[.] std::allocator<int>::allocator
1.79%	main	main	[.] std::vector<int, std::allocator<int> >::operator=
1.61%	main	main	[.] StringProcess
1.44%	main	main	[.] std::vector<int, std::allocator<int> >::~vector
1.43%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::get_allocator

Tip: See assembly instructions with percentage: perf annotate <symbol>

Samples: 143K of event 'cycles:u', Event count (approx.): 76376600377

Overhead	Command	Shared Object	Symbol
23.35%	main	main	[.] MD5Hash
9.54%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_impl_data::_M_copy_data
5.24%	main	main	[.] std::_Vector_base<segment, std::allocator<segment> >::_Vector_impl_data::_M_cop
3.69%	main	main	[.] PT::operator=
3.41%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_impl_data::_Vector_impl_d
2.32%	main	main	[.] std::vector<int, std::allocator<int> >::_M_move_assign
2.11%	main	main	[.] std::_Vector_base<segment, std::allocator<segment> >::_Vector_impl_data::_Vecto
1.75%	main	main	[.] std::allocator<int>::allocator
1.75%	main	main	[.] std::vector<int, std::allocator<int> >::operator=
1.51%	main	main	[.] StringProcess
1.42%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::get_allocator
1.39%	main	main	[.] std::vector<int, std::allocator<int> >::~~vector
1.38%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::~~_Vector_base
1.26%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_base
1.10%	main	main	[.] std::_do_alloc_on_move<std::allocator<int> >
1.08%	main	libstdc++.so.6.0.28	[.] std::_cxx11::basic_string<char, std::char_traits<char>, std::allocator<char> >
1.07%	main	main	[.] std::allocator<int>::~allocator
1.04%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_impl::_Vector_impl
0.98%	main	main	[.] std::vector<segment, std::allocator<segment> >::operator=
0.97%	main	main	[.] std::vector<segment, std::allocator<segment> >::_M_move_assign
0.96%	main	main	[.] std::vector<int, std::allocator<int> >::vector
0.93%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_M_deallocate
0.92%	main	main	[.] std::_alloc_on_move<std::allocator<int> >
0.89%	main	libc.so.6	[.] 0x000000000009c080
0.86%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_impl_data::_M_swap_data
0.82%	main	main	[.] __gnu_cxx::new_allocator<int>::~new_allocator
0.81%	main	main	[.] std::vector<segment, std::allocator<segment> >::~~vector
0.80%	main	main	[.] std::allocator<segment>::allocator
0.80%	main	main	[.] std::_Destroy<int*>
0.77%	main	main	[.] std::move<std::allocator<int>&>
0.77%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_M_get_Tp_allocator
0.77%	main	main	[.] std::_Vector_base<segment, std::allocator<segment> >::~~_Vector_base
0.71%	main	main	[.] std::_Destroy_aux<false>::__destroy<segment*>
0.67%	main	main	[.] std::_do_alloc_on_move<std::allocator<segment> >
0.64%	main	main	[.] std::allocator<segment>::~allocator
0.63%	main	main	[.] std::_Destroy<segment*>
0.62%	main	main	[.] std::_alloc_on_move<std::allocator<segment> >
0.61%	main	main	[.] std::__copy_move_backward<true, false, std::random_access_iterator_tag>::__copy
0.58%	main	main	[.] std::_Vector_base<segment, std::allocator<segment> >::get_allocator
0.54%	main	main	[.] std::_Vector_base<segment, std::allocator<segment> >::_Vector_impl_data::_M_swa
0.53%	main	main	[.] std::__copy_move<true, false, std::random_access_iterator_tag>::__copy_m<PT*, P
0.53%	main	main	[.] std::move<std::vector<int, std::allocator<int> >&&>
0.52%	main	main	[.] std::_Vector_base<segment, std::allocator<segment> >::_Vector_impl::_Vector_imp
0.52%	main	main	[.] std::vector<segment, std::allocator<segment> >::vector
0.52%	main	libc.so.6	[.] malloc
0.49%	main	main	[.] __gnu_cxx::new_allocator<segment>::~new_allocator

来个长图比对一下

Samples: 165K of event 'cycles:u', Event count (approx.): 88185461034

Overhead	Command	Shared Object	Symbol
19.18%	main	main	[.] MD5Hash_SIMD
8.45%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_impl_data::_M_copy_data
7.78%	main	main	[.] ROTATELEFT_SIMD
4.81%	main	main	[.] std::_Vector_base<segment, std::allocator<segment> >::_Vector_impl_data::_M_copy
3.17%	main	main	[.] PT::operator=
2.90%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_impl_data::_Vector_impl_da
2.09%	main	main	[.] std::vector<int, std::allocator<int> >::_M_move_assign
2.02%	main	main	[.] F_SIMD
1.98%	main	main	[.] G_SIMD
1.88%	main	main	[.] std::_Vector_base<segment, std::allocator<segment> >::_Vector_impl_data::_Vector
1.74%	main	main	[.] H_SIMD
1.58%	main	main	[.] std::vector<int, std::allocator<int> >::operator=
1.51%	main	main	[.] I_SIMD
1.50%	main	main	[.] std::allocator<int>::allocator
1.34%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::~~Vector_base
1.27%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::get_allocator
1.22%	main	main	[.] std::vector<int, std::allocator<int> >::~~vector
1.22%	main	main	[.] PrepareMessage
1.00%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_base
0.98%	main	main	[.] std::_do_alloc_on_move<std::allocator<int> >
0.96%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_impl::_Vector_impl
0.90%	main	main	[.] std::allocator<int>::~~allocator
0.88%	main	main	[.] std::vector<segment, std::allocator<segment> >::operator=
0.86%	main	main	[.] std::_Destroy<int*>
0.83%	main	main	[.] std::vector<int, std::allocator<int> >::vector
0.83%	main	main	[.] std::vector<segment, std::allocator<segment> >::_M_move_assign
0.83%	main	main	[.] std::_Vector_base<segment, std::allocator<segment> >::~~Vector_base
0.82%	main	main	[.] __gnu_cxx::new_allocator<int>::~~new_allocator
0.82%	main	main	[.] std::_alloc_on_move<std::allocator<int> >
0.78%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_M_deallocate
0.76%	main	libstdc++.so.6.0.28	[.] std::__cxx11::basic_string<char, std::char_traits<char>, std::allocator<char> >::
0.76%	main	main	[.] std::allocator<segment>::allocator
0.75%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_impl_data::_M_swap_data
0.73%	main	main	[.] std::_Destroy_aux<false>::_destroy<segment*>
0.72%	main	main	[.] std::vector<segment, std::allocator<segment> >::~~vector
0.70%	main	main	[.] std::move<std::allocator<int>&>
0.68%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_M_get_Tp_allocator
0.65%	main	main	[.] std::_Destroy<segment*>
0.65%	main	main	[.] std::_do_alloc_on_move<std::allocator<segment> >
0.61%	main	main	[.] std::allocator<segment>::~~allocator
0.59%	main	main	[.] std::_alloc_on_move<std::allocator<segment> >
0.52%	main	main	[.] std::move<std::vector<int, std::allocator<int> >&>
0.51%	main	main	[.] std::_Vector_base<segment, std::allocator<segment> >::get_allocator
0.51%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_M_get_Tp_allocator
0.50%	main	main	[.] std::_copy_move<true, false, std::random_access_iterator_tag>::_copy_m<PT*, PT
0.48%	main	main	[.] std::_copy_move_backward<true, false, std::random_access_iterator_tag>::_copy

Tip: To add Node.js USDT (User-Level Statically Defined Tracing): perf buildid-cache --add `which node`

豪德我们选中了内联函数ROTATELEFT_SIMD,并将其变为宏定义, 时间直接减少了1s

```
// ROTATELEFT的SIMD版本
inline bit32x4_t ROTATELEFT_SIMD(bit32x4_t x, int n) {
    return vorrq_u32(vshlq_n_u32(x, n), vshrq_n_u32(x, 32 - n));
}

#define ROTATELEFT_SIMD(num, n) (vorrq_u32(vshlq_n_u32((num), (n)),
vshrq_n_u32((num), 32 - (n))))
```

```
Guesses generated: 989/1/1
Guesses generated: 9997458
Guesses generated: 10097691
Guess time:30.8781seconds
Hash time:18.0347seconds
Train time:1.43003seconds
[ perf record: Woken up 58 times to write data ]
[ perf record: Captured and wrote 14.952 MB perf.data (322126 samples) ]
• [s2313211@master_ubss1 guess]$ perf report
```


Samples: 161K of event 'cycles:u', Event count (approx.): 86843961874

Overhead	Command	Shared Object	Symbol
24.95%	main	main	[.] MD5Hash_SIMD
8.32%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_impl_data::_M_copy_data
4.82%	main	main	[.] std::_Vector_base<segment, std::allocator<segment> >::_Vector_impl_data::_M_copy
3.22%	main	main	[.] PT::operator=
2.99%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_impl_data::_Vector_impl_da
2.38%	main	main	[.] F_SIMD
2.18%	main	main	[.] std::vector<int, std::allocator<int> >::_M_move_assign
1.99%	main	main	[.] std::_Vector_base<segment, std::allocator<segment> >::_Vector_impl_data::_Vector
1.92%	main	main	[.] G_SIMD
1.76%	main	main	[.] H_SIMD
1.73%	main	main	[.] I_SIMD
1.62%	main	main	[.] std::allocator<int>::allocator
1.61%	main	main	[.] std::vector<int, std::allocator<int> >::operator=
1.29%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::~~_Vector_base
1.28%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::get_allocator
1.22%	main	main	[.] PrepareMessage
1.21%	main	main	[.] std::vector<int, std::allocator<int> >::~~vector
1.00%	main	main	[.] std::_do_alloc_on_move<std::allocator<int> >
0.99%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_base
0.96%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_impl::_Vector_impl
0.95%	main	libstdc++.so.6.0.28	[.] std::__cxx11::basic_string<char, std::char_traits<char>, std::allocator<char> >::
0.94%	main	main	[.] std::allocator<int>::~~allocator
0.89%	main	main	[.] std::_alloc_on_move<std::allocator<int> >
0.89%	main	main	[.] std::vector<int, std::allocator<int> >::vector
0.86%	main	main	[.] std::vector<segment, std::allocator<segment> >::operator=
0.85%	main	main	[.] __gnu_cxx::new_allocator<int>::~~new_allocator
0.85%	main	main	[.] std::_Destroy<int*>
0.85%	main	main	[.] std::vector<segment, std::allocator<segment> >::_M_move_assign
0.79%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_M_deallocate
0.79%	main	main	[.] std::_Vector_base<segment, std::allocator<segment> >::~~_Vector_base
0.79%	main	libc.so.6	[.] 0x000000000009c080
0.76%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_impl_data::_M_swap_data
0.75%	main	main	[.] std::vector<segment, std::allocator<segment> >::~~vector
0.74%	main	main	[.] std::move<std::allocator<int>&>

ROTATELEFT也从高时间消失了

那我们不妨把这些也改为宏定义

```
// 添加SIMD版本的基本函数
inline bit32x4_t F_SIMD(bit32x4_t x, bit32x4_t y, bit32x4_t z) {
    return vorrq_u32(vandq_u32(x, y), vandq_u32(vmvnq_u32(x), z));
}

inline bit32x4_t G_SIMD(bit32x4_t x, bit32x4_t y, bit32x4_t z) {
    return vorrq_u32(vandq_u32(x, z), vandq_u32(y, vmvnq_u32(z)));
}

inline bit32x4_t H_SIMD(bit32x4_t x, bit32x4_t y, bit32x4_t z) {
    return veorq_u32(veorq_u32(x, y), z);
}

inline bit32x4_t I_SIMD(bit32x4_t x, bit32x4_t y, bit32x4_t z) {
    return veorq_u32(y, vorrq_u32(x, vmvnq_u32(z)));
}

#define F_SIMD(x, y, z) (vorrq_u32(vandq_u32((x), (y)), vandq_u32(vmvnq_u32((x)), (z))))
#define G_SIMD(x, y, z) (vorrq_u32(vandq_u32((x), (z)), vandq_u32((y), vmvnq_u32((z)))))
#define H_SIMD(x, y, z) (veorq_u32(veorq_u32((x), (y)), (z)))
#define I_SIMD(x, y, z) (veorq_u32((y), vorrq_u32((x), vmvnq_u32((z)))))
```



```

Guesses generated: 9595702
Guesses generated: 9696752
Guesses generated: 9796757
Guesses generated: 9897171
Guesses generated: 9997458
Guesses generated: 10097691
Guess time:30.9011seconds
Hash time:17.4475seconds
Train time:1.42324seconds
[ perf record: Woken up 56 times to write data ]
[ perf record: Captured and wrote 14.753 MB perf.data (318200 samples) ]

```

时间也来到了17.4s

问题 输出 调试控制台 终端 端口

perf - guess

Overhead	Command	Shared Object	Symbol
32.57%	main	main	[.] MD5Hash_SIMD
8.63%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_impl_data::_M_copy_data
4.90%	main	main	[.] std::_Vector_base<segment, std::allocator<segment> >::_Vector_impl_data::_M_copy
3.25%	main	main	[.] PT::operator=
3.05%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_impl_data::_Vector_impl_da
2.20%	main	main	[.] std::vector<int, std::allocator<int> >::_M_move_assign
2.00%	main	main	[.] std::_Vector_base<segment, std::allocator<segment> >::_Vector_impl_data::_Vector
1.64%	main	main	[.] std::allocator<int>::allocator
1.64%	main	main	[.] std::vector<int, std::allocator<int> >::operator=
1.31%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::~~_Vector_base
1.28%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::get_allocator
1.24%	main	main	[.] std::vector<int, std::allocator<int> >::~~vector
1.19%	main	main	[.] PrepareMessage
1.05%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_base
1.02%	main	main	[.] std::_do_alloc_on_move<std::allocator<int> >
1.00%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_impl::_Vector_impl
0.89%	main	main	[.] std::_Destroy<int*>
0.89%	main	main	[.] std::vector<segment, std::allocator<segment> >::operator=
0.88%	main	main	[.] std::vector<int, std::allocator<int> >::vector
0.88%	main	main	[.] std::allocator<int>::~~allocator
0.88%	main	main	[.] std::_alloc_on_move<std::allocator<int> >
0.87%	main	main	[.] std::vector<segment, std::allocator<segment> >::_M_move_assign
0.85%	main	libstdc++.so.6.0.28	[.] std::__cxx11::basic_string<char, std::char_traits<char>, std::allocator<char> >::
0.83%	main	main	[.] std::_Vector_base<segment, std::allocator<segment> >::~~_Vector_base
0.82%	main	main	[.] __gnu_cxx::new_allocator<int>::~~new_allocator
0.81%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_M_deallocate
0.75%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_impl_data::_M_swap_data
0.75%	main	main	[.] std::move<std::allocator<int>&>
0.74%	main	main	[.] std::allocator<segment>::allocator
0.71%	main	main	[.] std::_Destroy<segment*>
0.71%	main	main	[.] std::vector<segment, std::allocator<segment> >::~~vector
0.70%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_M_get_Tp_allocator
0.69%	main	main	[.] std::_Destroy_aux<false>::_destroy<segment*>
0.65%	main	main	[.] std::allocator<segment>::~~allocator
0.62%	main	main	[.] std::_do_alloc_on_move<std::allocator<segment> >
0.59%	main	main	[.] std::_alloc_on_move<std::allocator<segment> >
0.53%	main	main	[.] std::_Vector_base<segment, std::allocator<segment> >::get_allocator
0.53%	main	main	[.] std::move<std::vector<int, std::allocator<int> >&>
0.50%	main	main	[.] std::vector<segment, std::allocator<segment> >::vector
0.50%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_M_get_Tp_allocator

Tip: To count events in every 1000 msec: perf stat -i 1000

内联函数的时间消耗已经被消除了但是主函数的时间增加了，我们可以看出内联函数虽然可以减少主函数的时间，但自身会增加运行时间

我们观察就可以发现凹其他的时间已经没用了，只能去凹主函数

Percent	Command	Shared Object	Symbol
4.05	1d8: cmp x0, #0x0		
0.00	↓ b.lt 1f0		
0.06	strb wzr, [x1]		
0.01	add x1, x1, #0x1		
0.01	sub x0, x0, #0x1		
0.02	↑ b 1d8		
0.02	1f0: str x2, [sp, #20072]		

看到这个没有，百分之4.05，优化他！

```
memset(paddedMessage + length + 1, 0, paddingBytes - 1); // 填充0字节是这个语句
```

```
static Byte zero_buffer[MAX_BUFFER_SIZE] = {0}; // 加上一个静态缓冲区  
memcpy(output + input_length + 1, zero_buffer, padding_bytes - 1); // 用拷贝复制信息
```

```
-  
Guesses generated: 9490827  
Guesses generated: 9595702  
Guesses generated: 9696752  
Guesses generated: 9796757  
Guesses generated: 9897171  
Guesses generated: 9997458  
Guesses generated: 10097691  
Guess time:28.6479seconds  
Hash time:16.6006seconds  
Train time:1.26677seconds  
[ perf record: Woken up 46 times to write data ]  
[ perf record: Captured and wrote 11.763 MB perf.data (252457 samples) ]  
  
Guesses generated: 9595702  
Guesses generated: 9696752  
Guesses generated: 9796757  
Guesses generated: 9897171  
Guesses generated: 9997458  
Guesses generated: 10097691  
Guess time:28.5686seconds  
Hash time:16.555seconds  
Train time:1.29609seconds  
[ perf record: Woken up 45 times to write data ]  
[ perf record: Captured and wrote 11.741 MB perf.data (252128 samples) ]  
[ 252128 samples collected, 11.741 MB written ]
```

16.6s了已经

Percent		and	x0, x0, #0xfffffffffffffc0
0.11		str	x0, [sp, #20080]
		ldr	x0, [sp, #20080]
		lsl	x19, x0, #2
		mov	x0, x19
		→ bl	_init
		mov	x2, x0
		mov	x1, x2
		sub	x0, x19, #0x1
4.00	1d8:	cmp	x0, #0x0
		↓ b.lt	1f0
0.05		strb	wzr, [x1]
		add	x1, x1, #0x1
0.01		sub	x0, x0, #0x1
		↑ b	1d8
0.03	1f0:	str	x2, [sp, #20072]
		sub	x0, sp, #0x120
		add	x0, x0, #0x2c0
0.12		ldr	x1, [sp, #20072]
		str	x1, [x0]
		sub	x0, sp, #0x120
		add	x0, x0, #0x2c0
0.02		ldr	x0, [x0]
		add	x1, sp, #0x190
		mov	x2, x1
		mov	x1, x0

性能瓶颈的汇编代码，我们再定位一下代码位置

```
Byte* buffer = new Byte[4 * max_padded_length](); // 括号导致内存清零
```

```
Byte* buffer = new Byte[4 * max_padded_length]; // 只初始化必要的部分
```

```
-  
Guesses generated: 9997458  
Guesses generated: 10097691  
Guess time:28.7946seconds  
Hash time:15.6209seconds  
Train time:1.31909seconds  
[ perf record: Woken up 45 times to write data ]  
[ perf record: Captured and wrote 11.554 MB perf.data (249221 samples) ]  
○ [s2313211@master ubss1 guess]$ █
```

15.5s

Samples: 124K of event 'cycles:u', 2750 Hz, Event count (approx.): 84619676314

```

0.02      ldrb  w0, [x0]
          lsl  w0, w0, #16
          orr  w1, w1, w0
0.12      ldr   x0, [sp, #20024]
          add  x0, x0, #0x3
0.04      ldrb  w0, [x0]
0.01      lsl  w0, w0, #24
          orr  w0, w1, w0
0.17      str   w0, [sp, #308]
0.01      add  x0, sp, #0x128
0.03      str   x0, [sp, #19920]
0.16      ldr   x0, [sp, #19920]
0.17      ldr   q0, [x0]
0.03      nop
1.31      add  x0, sp, #0x4, lsl #12
          ldrrw x0, [x0, #3708]
          lsl  x0, x0, #4
          add  x1, sp, #0x20
0.00      str   q0, [x1, x0]
          add  x0, sp, #0x4, lsl #12
0.01      ldr   w0, [x0, #3708]
          add  w0, w0, #0x1
          add  x1, sp, #0x4, lsl #12
0.10      str   w0, [x1, #3708]
          ↑ b 394
0.00      5d4: ldr   q0, [sp, #20160]
0.01      str   q0, [sp, #20000]
          ldr   q0, [sp, #20144]
          str   q0, [sp, #19984]
          ldr   q0, [sp, #20128]
0.01      str   q0, [sp, #19968]
          ldr   q0, [sp, #20112]
          str   q0, [sp, #19952]
          ldr   q0, [sp, #19984]
0.00      str   q0, [sp, #19966]
          ldr   q0, [sp, #19968]
          .
          .
          .

```

定位代码

```

// 更高效的数据加载
for (int j = 0; j < 16; j++) {
    // 一次性加载4个值
    uint32_t values[4];
    // 展开的数据加载循环
    Byte* block_ptr0 = paddedMessages[0] + block * 64 + j * 4;
    Byte* block_ptr1 = paddedMessages[1] + block * 64 + j * 4;
    Byte* block_ptr2 = paddedMessages[2] + block * 64 + j * 4;
    Byte* block_ptr3 = paddedMessages[3] + block * 64 + j * 4;

    values[0] = (uint32_t)block_ptr0[0] | ((uint32_t)block_ptr0[1] << 8) |
((uint32_t)block_ptr0[2] << 16) | ((uint32_t)block_ptr0[3] << 24);
    values[1] = (uint32_t)block_ptr1[0] | ((uint32_t)block_ptr1[1] << 8) |
((uint32_t)block_ptr1[2] << 16) | ((uint32_t)block_ptr1[3] << 24);
    values[2] = (uint32_t)block_ptr2[0] | ((uint32_t)block_ptr2[1] << 8) |
((uint32_t)block_ptr2[2] << 16) | ((uint32_t)block_ptr2[3] << 24);
    values[3] = (uint32_t)block_ptr3[0] | ((uint32_t)block_ptr3[1] << 8) |
((uint32_t)block_ptr3[2] << 16) | ((uint32_t)block_ptr3[3] << 24);
    // 使用NEON指令直接加载数据
    M[j] = vld1q_u32(values);
}

```

```

static const size_t BLOCK_OFFSETS[16] = {0, 4, 8, 12, 16, 20, 24, 28, 32, 36, 40,
44, 48, 52, 56, 60};
//加上静态变量避免重复计算

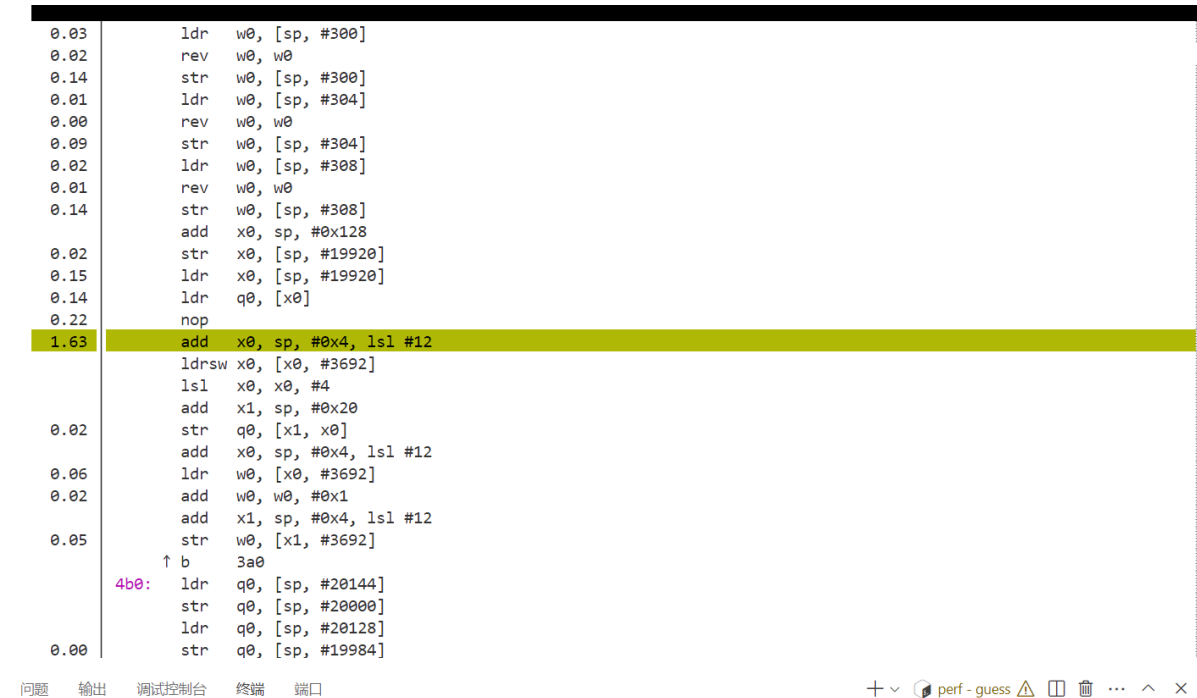
```

```

Guesses generated: 9796757
Guesses generated: 9897171
Guesses generated: 9997458
Guesses generated: 10097691
Guess time:28.9344seconds
Hash time:15.3386seconds
Train time:1.27717seconds
[ perf record: Woken up 43 times to write data ]
[ perf record: Captured and wrote 11.504 MB perf.data (247907 samples) ]
[s2313211@master_ubss1 guess]$

```

似乎没有优化，现在最多的变为了这个



```

0.03    ldr    w0, [sp, #300]
0.02    rev    w0, w0
0.14    str    w0, [sp, #300]
0.01    ldr    w0, [sp, #304]
0.00    rev    w0, w0
0.09    str    w0, [sp, #304]
0.02    ldr    w0, [sp, #308]
0.01    rev    w0, w0
0.14    str    w0, [sp, #308]
        add    x0, sp, #0x128
0.02    str    x0, [sp, #19920]
0.15    ldr    x0, [sp, #19920]
0.14    ldr    q0, [x0]
0.22    nop
1.63    add    x0, sp, #0x4, lsl #12
        ldrrw x0, [x0, #3692]
        lsl    x0, x0, #4
        add    x1, sp, #0x20
0.02    str    q0, [x1, x0]
        add    x0, sp, #0x4, lsl #12
0.06    ldr    w0, [x0, #3692]
0.02    add    w0, w0, #0x1
        add    x1, sp, #0x4, lsl #12
0.05    str    w0, [x1, #3692]
        ↑ b
4b0:    ldr    q0, [sp, #20144]
        str    q0, [sp, #20000]
        ldr    q0, [sp, #20128]
0.00    str    q0, [sp, #19984]

```

Samples: 122K of event 'cycles:u', Event count (approx.): 84194169277

Overhead	Command	Shared Object	Symbol
31.00%	main	main	[.] MD5Hash_SIMD
8.87%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_impl_data::_M_copy_data
5.42%	main	main	[.] std::_Vector_base<segment, std::allocator<segment> >::_Vector_impl_data::_M_copy
3.29%	main	main	[.] PT::operator=
2.97%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_impl_data::_Vector_impl_da
1.94%	main	main	[.] std::vector<int, std::allocator<int> >::_M_move_assign
1.86%	main	main	[.] std::allocator<int>::allocator
1.65%	main	main	[.] std::_Vector_base<segment, std::allocator<segment> >::_Vector_impl_data::_Vector
1.52%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::get_allocator
1.47%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::~~Vector_base
1.46%	main	main	[.] std::vector<int, std::allocator<int> >::operator=
1.24%	main	main	[.] std::vector<int, std::allocator<int> >::~~vector
1.13%	main	main	[.] PrepareMessage
1.12%	main	main	[.] std::__do_alloc_on_move<std::allocator<int> >
1.11%	main	main	[.] std::_Destroy<int*>
1.08%	main	main	[.] std::vector<segment, std::allocator<segment> >::_M_move_assign
1.02%	main	main	[.] std::allocator<int>::~~allocator
1.00%	main	main	[.] std::__alloc_on_move<std::allocator<int> >
0.89%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_base
0.89%	main	main	[.] __gnu_cxx::new_allocator<int>::~~new_allocator
0.89%	main	main	[.] std::allocator<segment>::allocator
0.89%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_impl_data::_M_swap_data
0.88%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_Vector_impl::_Vector_impl
0.88%	main	main	[.] std::vector<segment, std::allocator<segment> >::~~vector
0.85%	main	main	[.] std::vector<segment, std::allocator<segment> >::operator=
0.82%	main	main	[.] std::move<std::allocator<int>&>
0.79%	main	main	[.] std::_Vector_base<int, std::allocator<int> >::_M_deallocate
0.77%	main	main	[.] std::_Destroy_aux<false>::_destroy<segment*>
0.74%	main	main	[.] std::_Vector_base<segment, std::allocator<segment> >::get_allocator
0.73%	main	libstdc++.so.6.0.28	[.] std::__cxx11::basic_string<char, std::char_traits<char>, std::allocator<char> >::
0.68%	main	main	[.] std::_Destroy<segment*>

```
110 guesses generated: 9097171
111 Guesses generated: 9997458
112 Guesses generated: 10097691
113 Guess time:20.0164seconds
114 Hash time:10.9214seconds
115 Train time:0.8663seconds
116 

---


117 Authorized users only. All activities may be monitored and reported.
118
```